Map-reduce Computing for Introductory Students using WebMapReduce

CSInParallel Project

June 10, 2016

CONTENTS

Last Updated: 2016-06-10

The first section below is preliminary reading for any of the following three sections. It describes a map-reduce system, or framework, using distrubted computers in a cluster to carry out analysis of massive amounts of data concurently, or "in parallel", by having processes on each computer work on small portions of a much larger dataset. The next section, *Using WebMapReduce (WMR)*, walks you through a simple, basic example of counting words in text files. The third section, entitled *Counting words with WebMapReduce (WMR): adding efficiency*, shows a slightly revised version of the basic word counting that is more efficient than basic example, with some small modifications to the code and the files it works on. The fourth section provides further example data and analyses that shows the range of what you can do with map-reduce systems.

USING PARALLELISM TO ANALYZE VERY LARGE FILES: GOOGLE'S MAP REDUCE PARADIGM

1.1 Introduction

Up to now in this course, you have been writing programs that run *sequentially* by executing one step at a time until they complete (unless you mistakenly wrote one with an infinite loop!). Sequential execution is the standard basic mode of operation on all computers. However, when we need to process large amounts of data that might take a long time to complete, we can use a technique known as *parallel computing*, in which we define portions of our computation to run at the same time by using multiple processors on a single machine or multiple computers at the same time.

In a lab and homework in this course you will use a system called *MapReduce* that was designed to harness the power of many computers together in a *cluster* to execute in parallel to process large amounts of data. You will be writing programs designed to run in parallel, or concurrently (at the same time), each one working on a portion of the data. By doing this, your parallel program will be faster than a corresponding sequential version using one computer to process all of the data. We call this kind of program one that uses *data parallelism*, because the data is split and processed in parallel.

MapReduce is a strategy for solving problems using two stages for processing data that includes a sort action between those two stages. This problem-solving strategy has been used for decades in *functional programming* languages such as LISP or Scheme. More recently Google has adapted the map-reduce programming model (Dean and Ghemawat,2004) to function efficiently on large clusters of computers to process vast amounts of data–for example, Google's selection of the entire web.

The Apache Foundation provides an open-source implementation of MapReduce for clusters called Hadoop, which has primarily been implemented by Yahoo!. Student researchers at St. Olaf College have created a user interface called *WebMapReduce (WMR)* that uses Hadoop to make map-reduce programming convenient enough for students in computer science courses to use.

1.2 Description of the MapReduce programming model

In map-reduce programming, a programmer provides two functions, called the *mapper* and the *reducer*, for carrying out a sequence of two computational stages on potentially vast quantities of data. A series of identical 'map' functions can be run on a large amount of input data in parallel, as shown in Figure 1. The results from these mapper processes, spread across many computers, are then sent to reduce functions, also spread across many computers. The most important concepts to remember are these:

- The mapper function that you design is applied to each line of data, and breaks up that data into labelled units of interest, called *key-value pairs*.

- The reducer function that you design is then applied to all the key-value pairs *that share the same key*, allowing some kind of consolidation of those pairs.
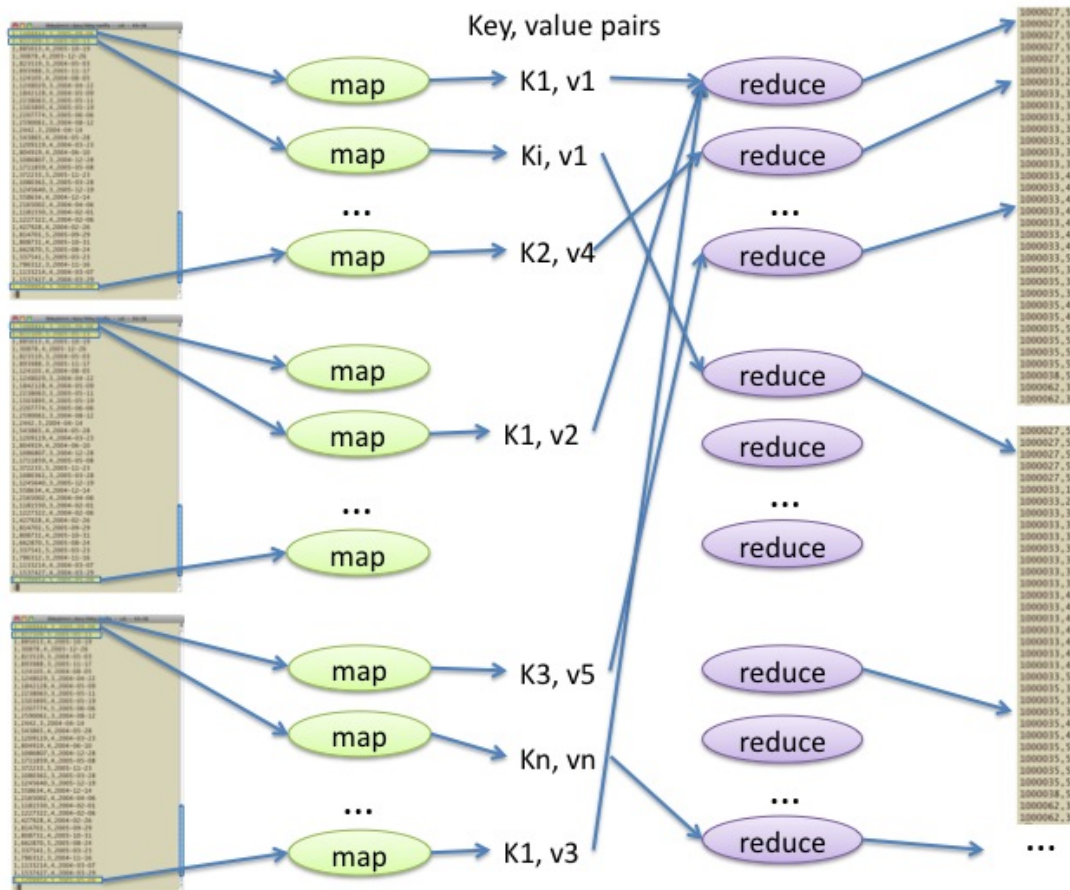

Figure 1.1: Figure 1: The concept behind how map functions can run in parallel and pass their results to reduce functions, whose results are output in sorted order by the keys created by the reduce function.

In a *map-reduce system*, which is made of of many computers working at the same time, some computers are assigned mapper tasks, some shuffle the data from mappers and hand it over to reducers, and some computers handle the reducer tasks. Between the mapper and reducer stages, a map-reduce system automatically reorganizes the intermediate key-value pairs, so that each call of the reducer function can receive a complete set of key-value pairs *for a particular key*, and so that the reducer function is called for every key in sorted order. We will refer to this reorganization of key-value pairs between the mapper and reducer stages as *shuffling*. Figure 2 illustrates the three steps of mapping, shuffling, and reducing.

Before the mapper stage, a map-reduce system such as Hadoop breaks the entire input data up into portions, known as *splits*; each split is assigned to one computer in a cluster, which then applies the mapper function to each line of that split. Thus,multiple computers can process different splits of data simultaneously. (With, say, 2000 computers or *nodes* in a large-scale commercial cluster, quadrillions of bytes (*petabytes*) of data might be processed in a few hours that might otherwise require years of computing time on a single computer.) Likewise, reducers may be run on multiple computers in order to speed up the performance of that stage.
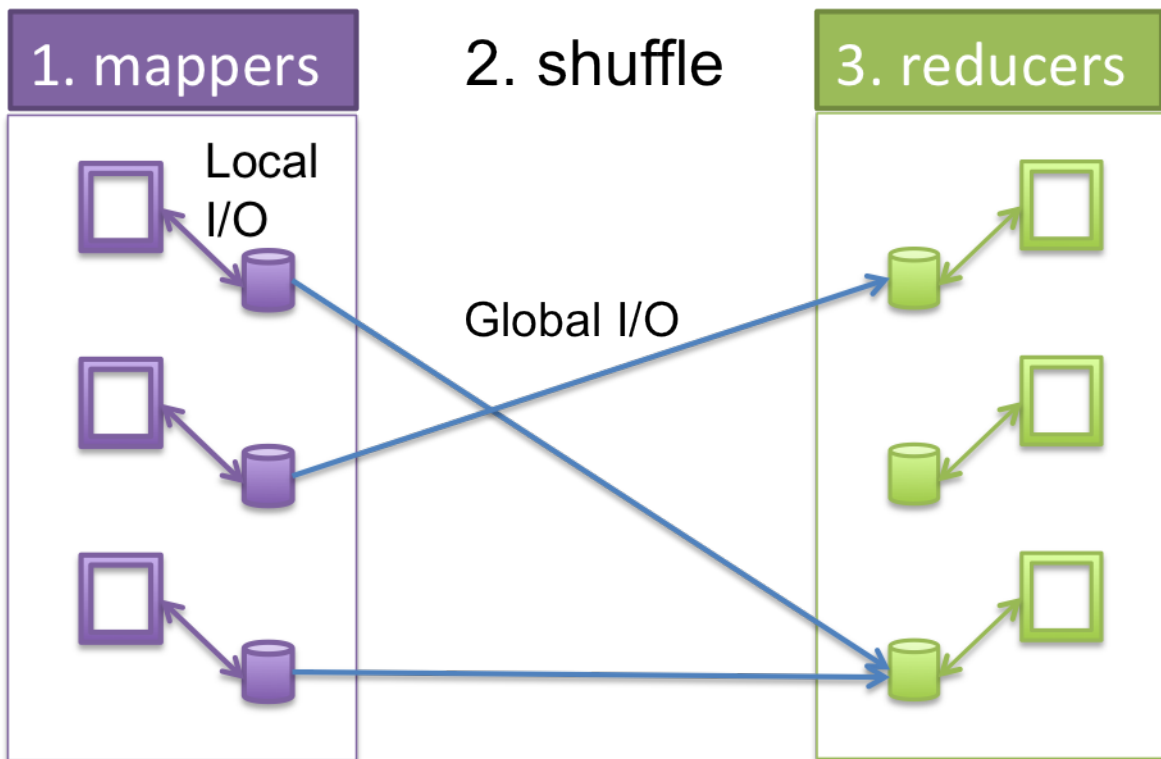
Figure 1.2: Figure 2: How each computer in a cluster breaks up the work and runs mappers locally, then shuffles the key-value pair results by key and sends the results for each key to other computers who run reducers.

Parallel computing is the practice of using multiple computations at the same time in order to improve the performance of those computations. The map-reduce programming model is an example of two varieties of parallel computing:

1. *Data parallelism*, in which multiple portions of data are processed simultaneously on multiple processors (CPUs). This occurs when multiple splits of data are handled on different computers in a cluster.

2. *Pipelining*, in which data is processed in a sequence of stages, like an assembly line. The mapper and reducer stages represent a two-stage pipeline. If shuffling is considered as its own stage, the pipeline would have three stages. Pipelining is an example of *task parallelism*, in which different computational tasks take place at the same time. In the case of the map-reduce stages, mapping could overlap with shuffling to some extent, by having mappers stream their output to shuffle processes, which would prepare it for reducers while the mappers are generating results. Thus, different computers could handle each of these tasks.

Note: Hadoop actually carries out the three stages of mapping, shuffling, and reducing *sequentially* (one stage after the other), instead of using task parallelism. That is, all of the mapping occurs before any of the shuffling begins, and all of the shuffling is completed before any of the reducing begins. (See below for reasons why.) Thus, Hadoop's implementation of map-reduce doesn't literally qualify as pipeline parallelism, because multiple stages do not take place at the same time. However, true pipeline parallelism *does* take place within our testing program used in the `WebMapReduce` interface, called `wmrtest`, which is useful for testing and debugging mapper and reducer functions with small data. In general, solving problems using pipeline (assembly line) thinking creates opportunities for using parallelism to improve performance.

Fault tolerance. Large (e.g., 2000-node) clusters offer the potential for great performance speedup, but breakdowns are inevitable when using so many computers. In order to avoid losing the results of enormous computations because of breakdowns, map-reduce systems such as Hadoop are *fault tolerant*, i.e., designed to recover from a significant amount of computer failure. Here are some fault-tolerance strategies used in Hadoop:

- All data is *replicated* on multiple computers. Thus,if one computer fails, its data is readily available on other computers.

- If a computer running a mapper function fails, that mapper operation can thus be restarted on another computer that has the same data (after discarding the partial results (key-value pairs) from incomplete mapper calls on that failed computer).

- If all mappers have completed, but a computer fails during shuffling, then any lost mapper results can be regenerated on another computer, and shuffling can resume using non-failed computers.

- Shuffling results are also replicated, so if a computer running reducers fails, those reducers can be rerun on another computer.

Note: Hadoop's fault tolerance features make it a good use for the *selkie* cluster at Macalester, which uses the many computers in two large labs in the MSCS Department in Olin-Rice. Occasionally, these are sometimes unfortunately rebooted by users. These occasional failures of machines in the cluster can be compensated for by Hadoop. However, when many of these computers are rebooted at about the same time, all of the copies of some data may become unavailable, leading to Hadoop failures.

USING WEBMAPREDUCE (WMR)

2.1 Introduction

For this activity, you need to have read the accompanying background reading in the previous section entitled *Using Parallelism to Analyze Very Large Files: Google's MapReduce Paradigm.*

In this lab you will use a web-based program called *WebMapReduce* (**WMR**) that enables you to formulate map-reduce computations and run them on an underlying *Hadoop* map-reduce system running on a cluster of computers.

You will use WebMapReduce on a cluster of computers at Saint Olaf College. You should access WebMapReduce now and register for a login by going to this URL on your web browser:

http://cumulus.cs.stolaf.edu/wmr/

Choose the link at the very upper right of this page that says 'Register'. Use your email address as your login name, and provide the other information asked for. You choose your own password, so that you can remember it and have control of your account.

Note: To complete your registration, you will need a registration key from your instructor.

For later reference, you may want to check this documentation for WMR.

For this activity, you should be able to follow along with the instructions below and determine how to use WMR.

2.2 An example of map-reduce computing with WMR: counting words

To program with map-reduce, you must first decide how to use mappers and reducers in order to accomplish your computational goal. The mapper function should take a line of input and decompose it somehow into key-value pairs; then the reducer should somehow condense or analyze all the key-value pairs having a *common key*, and produce a desired result.

The following example is small to illustrate how the process works. In realistic applications, which you will try later in this activity and in homework, the input data is much larger (several to hundreds of Gigabytes) and is stored in the Hadoop system. You will go through the following exercise first to ensure that the code is working and that you understand how it works. Then you can move on to bigger files. This is the process you should go through when doing this kind of coding: work on small amounts of data first to ensure correctness of your code, then try larger amounts of data.

As an example, consider the problem of counting how frequently each word appears in a collection of text data. For example, if the input data in a file is:

```
One fish, Two fish,
Red Fish, Blue fish.
Blue Fish, Two Fish.
```

then the output should be:

```
Blue 2
One 1
Red 1
Two 2
fish, 2
Fish, 2
fish. 1
Fish. 1
```

As this output indicates, we did not make any attempt to trim punctuation characters in this first example. Nor did we consider that some words in our text might be capatalized and some may not. We will also not do so as we practice using WebMapReduce with the initial functions described below. However, you can consider adding punctuation removal and lowercase conversion to your mapper code as you work through the example.

Note: The WebMapReduce system will sort the words according to the ASCII codes of the characters within words.

What follows is a plan for the mapper and reducer functions.

2.2.1 Map-reduce plan

In WMR, mapper functions work simultaneously on lines of input from files, where a line ends with a newline charater. The mapper will produce one key-value pair (w, *count*) for each and every word encountered in the input line that it is working on.

Thus, on the above input, three mappers working together, one on each line, could emit the following combined output:

```
One 1
fish, 1
fish, 1
Two 1

Red 1
Fish, 1
Blue 1
fish. 1

Blue 1
Fish, 1
Two 1
Fish. 1
```

The reducers will compute the sum of all the *count* values for a given word w, then produce the key-value pair (w, *sum*).

In this example, we can envision a reducer for each distinct word found by the three mappers, where the reducer gets a list of single counts per occurance of the word that a mapper found, loking like this:

```
One     [1]
fish,   [1, 1]
Two     [1,1]
Red     [1]
```

```
Fish,  [1,1]
Blue   [1,1]
fish.  [1]
Fish.  [1]
```

Each reducer works on one of the pairs of data shown above, and the system handles creating words with the lists of counts as shown above.

2.3 One System, Many Languages

In map-reduce framework systems in general and in WMR specifically, you can use one of several programming languages to code your mapper and reducer functions. The following table contains links to solutions in several languages for the word-count solution we describe below.

Lnaguage	Mapper function code	Reducer function code
Python	`wcmapper.py`	`wcreducer.py`
C++	`wcmapper.cpp`	`wcreducer.cpp`
C	`wcmapper.c`	`wcreducer.c`
Java	`wcmapper.java`	`wcreducer.java`

2.4 The mapper function

Each mapper process is receiving a line from a file as its key initially when the process starts (the value is empty, or null). You write one mapper function to be executed by those prcesses on any given line from any particular file. Our goal is to have the mapper output a new (key, value) containing a word found and the number one, as shown for the three-mapper example above.

Here is psedocode for what a typical mapper might accomplish:

```
# key is a single line from a file.
#
# value is empty in this case, since this is the first mapper function
# we are applying.
#
function mapper(key, value)
  1) Take the key argument to this function, which is the line of text,
     and split it on whitespace

  2) For every word resulting from the split key line:

      'emit' a pair (word, "1") to the system for the reducers to handle
```

Here is a Python3 mapper function for accomplishing this task using the WMR system. We include the feature of stripping away puncuation characters from the input and converting each word found to lowercase.

```python
def mapper(key, value):
    words=key.split()
    for word in words:
        Wmr.emit(word, '1')
```

This code is available for download in the table above, as are versions in other languages. Note that in each language you will need to know how to specify the (key, value) pair to emit to the system for the reducers to process. We see this for Python in line 4 above.

2.5 The reducer function

In the system, there will be reducer processes to handle each word 'key' that was emitted by various mappers. You write reducer code as if your reducer function is getting one word key and a container of counts, where each count came from a different mapper that was working on a different line of a file or files. In this simplest example, each count is a '1', each of which will be summed together by the reducer handling the particular word as a key.

Pseudocode for a reducer for this problem looks like this:

```
# key is a single word, values is a 'container' of counts that were
# gathered by the system from every mapper
#
function reducer(key, values)

  set a sum accumulator to zero

  for each count in values
    accumulate the sum by adding count to it

  'emit' the (key, sum) pair
```

A reducer function for solving the word-count problem in Python is

```
1  def reducer(key, iter):
2      sum = 0
3      for s in iter:
4          sum = sum + int(s)
5      Wmr.emit(key, str(sum))
```

This code is also available in the table above containing versions in several languages.

The function `reducer()` is called once for each distinct key that appears among the key-value pairs emitted by the mapper, and that call processes all of the key-value pairs that use that key. On line 1, the two parameters that are arguments of `reducer()` are that one distinct `key` and a Python3 *iterator* (similar to a list, but not quite) called `values`, which provides access to all the values for that key. Iterators in Python3 are designed for `for` loops- note in line 3 that we can simply ask for each value one at a time from the set of values held by the iterator.

Rationale: WMR `reducer()` functions use iterators instead of lists because the number of values may be very large in the case of large data. For example, there would be billions of occurrences of the word "the" if our data consisted of all pages on the web. When there are a lot of key-value pairs, it is more efficient to dispense pairs one at a time through an iterator than to create a gigantic complete list and hold that list in main memory; also, an enormous list may overfill main memory.

The `reducer()` function adds up all the counts that appear in key-value pairs for the `key` that appears as `reducer()`'s first argument (recall these come from separate mappers). Each count provided by the iterator `values` is a string, so in line 4 we must first convert it to an integer before adding it to the running total `sum`.

The method `Wmr.emit()` is used to produce key-value pairs as output from the mapper. This time, only one pair is emitted, consisting of the word being counted and `sum`, which holds the number of times that word appeared in *all* of the original data.

2.6 Running the example code on WebMapReduce

To run WMR with this combination of data, mapper, and reducer, carry out the following steps.

- In a browser, visit the WMR site at (if you don't already have it open from registering):

    http://cumulus.cs.stolaf.edu/wmr/

- After you have registered, you can use your email address and password to login. After successfully logging in, you are taken to the WMR page where you can complete your work.

- Enter a job name (perhaps involving your username, for uniqueness; avoid spaces in the job name and make sure that it is more than 4 characters long).

- Choose the language that you wish to try.

- For now, you can leave the number of map tasks and reduce tasks blank. This will let the system decide this for itself. You can also leave the default choice of sorting alphabetically.

- Enter the input data, e.g., the fish lines above. You can use the *Direct Input* option and enter that data in the text box provided.

- Enter the mapper. It's probably best to use se the ''Upload'' option and navigate to a file that contains the mapper, which you have entered using an editor (this is more convenient for repeated runs). Or you can use the file we provided in a table of links above.

  Beware: cutting and pasting your code from a pdf file or a web page or typing it into the ''direct' entry box for Python code is a bit problematic, because the needed tabs in the code might not be preserved (although using spaces should work). Check that the appropriate radio button is clicked to indicate the source option you're actually using.

- Also enter the reducer. Again, it's easier to use the file provided with a link in the table above.

- Click on the submit button.

A page should appear indicating that the job started successfully. This page will refresh itself as it is working on the job to show you progress.

Once the job runs to completion, you should see a Job Complete page. This page will include your output. If you used the fish input, your output should match the illustration above, except that the punctuation should also be taken care of.

If something doesn't work as described here, the following section may help with troubleshooting. *Read it next in any case so that you know what you can do when you work on your own new examples.*

2.7 Using WMR and its test mode

Here is some information about developing WMR map-reduce programs,and what to do if something goes wrong with your WMR job.

- First, a reminder:

  - At present, the WMR interface does not automatically reset radio buttons for you when you upload a file or use *Distributed FileSystem* data generated from a prior map-reduce run. *Always check to see that the radio buttons select the data, mapper, and reducer resources you intend.*

- You can test your mapper alone without using your reducer by using the *identity reducer*, which simply emits the same key-value pairs that it receives. Here is an implementation of the identity reducer for Python:

```
1  def reducer(key, iter):
2      for s in iter:
3          Wmr.emit(key, s)
```

As an example, if you use the word-count mapper with the identity reducer, then the "fish" data above should produce the following output:

```
Blue 1
Blue  1
fish, 1
fish, 1
fish. 1
Fish, 1
Fish, 1
Fish. 1
One 1
Red 1
Two 1
Two 1
```

Observe that the output is sorted, due to the shuffling step. However, this does show all the key-value pairs that result from the word-count mapper.

- Likewise, you can test your reducer alone without using your mapper by substituting the `identity mapper`, which simply copies key-value pairs from lines of input data. Here is an implementation of the identity mapper in Python:

```
1  def mapper(key, value):
2      Wmr.emit(key, value)
```

Here are identity mappers and reducers for some languages.

Lnaguage	Mapper function code	Reducer function code
Python	idmapper.py	idreducer.py
C++	idmapper.cpp	idreducer.cpp
Java	idmapper.java	idreducer.java

For example, you could enter a small amount of input data that you expect your mapper to produce, such as the `TAB`-separated key-value pairs listed above from using the identity reducer. If you then use the identity mapper `idmapper.py` with the word-count reducer `wcreducer.py` you should get the following output, which we would expect from each stage working:

```
Blue 2
fish, 2
fish. 1
Fish, 2
Fish. 1
One 1
Red 1
Two 2
```

Note: Use a `TAB` character to separate the key and value in the input lines above. To keep a test case around, it is easier to enter your data in an editor, then cut and paste to enter that data in the text box. Alternatively, you can"Upload" a file that contains the data.

- Unfortunately, the current WMR system does *not* provide very useful error messages in many cases. For example, if there's an error in your Python code, no clue about that error can be passed along in the current system.

- In order to test or debug a mapper and reducer, you can use the `Test` Button at the bottom of the WMR web page. The job output from this choice shows you what both the mapper and reducer emitted, which can be helpful for debugging your code.

Note: Do not use `Test` for large data, but only to debug your mappers and reducers. This option does *not* use cluster computing, so it cannot handle large data.

2.8 Next Steps

1. In WMR, you can choose to use your own input data files. Try choosing to 'browse' for a file, and using this `mobydick.txt` as file input.

2. You have likely noticed that capitalized words are treated separately from lowercase words. Change your mapper to also convert each word to lowercase before checking whether it is in the dictionary.

3. Also remove punctuation from each word after splitting the line (so all the 'fish' variations get counted together).

4. There are a large number of files of books from Project Gutenberg available on the Hadoop system underlying WebMapReduce. First start by trying this book as an input file by choosing 'Cluster Path' as Input in WMR and typing one of these into the text box:

   /shared/gutenberg/WarAndPeace.txt

   /shared/gutenberg/CompleteShakespeare.txt

   /shared/gutenberg/AnnaKarenina.txt

   These books have many lines of text with 'newline' chacters at the end of each line. Each of you mapper functions works on one line. Try one of these.

5. Next, you should try a collection of many books, each of which has no newline characters in them. In this case, each mapper 'task' in Hadoop will work on one whole book (your dictionary of words per mapper will be for the whole book, and the mappers will be running on many books at one time). In the Hadoop file system inside WMR we have these datasets available for this:

'Cluster path' to enter in WMR	Number of books
/shared/gutenberg/all_nonl/group10	2018
/shared/gutenberg/all_nonl/group11	294
/shared/gutenberg/all_nonl/group6	830
/shared/gutenberg/all_nonl/group8	541

   While using many books, it will be useful for you to experiment with the different datasets so that you can get a sense for how much a system like Hadoop can process.

   To do this, it will also be useful for you to save your configuration so that you can use it again with a different number of reducer tasks. Once you have entered your mapper and reducer code, picked the Python3 language, and given your job a descriptive name, choose the *'Save'* button at the bottom of the WMR panel. This will now be a *'Saved Configuration'* that you can retrieve using the link on the left in the WMR page.

   Try using the smallest set first (group11). Do not enter anything in the map tasks box- notice that the system will choose the same number of mappers as the number of books (this will show up once you submit the job). Also do not enter anything for the number of reduce tasks. With that many books, when the job completes you will see there are many pages of output, and some interesting 'words'. For the 294 books in group11, note how you obtain several pages of results. You will also notice that the stripping of punctuation isn't perfect. If you wish to try improving this you could, but it is not necessary.

2.9 Additional Notes

It is possible that input data files to mappers may be treated differently than as described in the above example. For example, a mapper function might be used as a second pass by operating on the reducer results from a previous map-reduce cycle. Or the data may be formatted differently than a text file from a novel or poem. These notes pertain to those cases.

In WMR, each line of data is treated as a key-value pair of strings, where the key consists of all characters before the first `TAB` character in that line, and the value consists of all characters after that first `TAB` character. Each call of `mapper()` operates on one line and that function's two arguments are the key and value from that line.

If there are multiple `TAB` characters in a line, all `TAB` characters after the first remain as part of the `value` argument of `mapper()`.

If there are *no* `TAB` characters in a data line (as is the case with all of our fish lines above), an empty string is created for the value and the entire line's content is treated as the key. This is why the key is split in the mapper shown above.

COUNTING WORDS WITH WEBMAPREDUCE (WMR): ADDING EFFICIENCY

3.1 Introduction

For this activity, you need to have read the accompanying background reading in the first section entitled *Using Parallelism to Analyze Very Large Files: Google's MapReduce Paradigm*. Also, you should have gone through the previous section where you learned to use WMR on a simple word-counting example.

WMR has several languages as options. In this case we will demonstrate an improvement that can be made easily if you are using the Python language, because it has dictionaries, or hash maps, as a built-in data type.

3.2 An improved word-count example

As in the previous section, we will start with a small example as an illustration. In this case we will describe an improvement in which mappers do a bit more work by keeping counts of words it has encountered and then emitting each word and its total count to the system for the reducer processes to handle. In a map-reduce system, it turns out to be useful to let the mappers do a fair amount of work, such as processing a whole book, since this is a reasonable task for a single process.

As an example, consider the problem of counting how frequently each word appears in a collection of data. For example, if the input data in a file is:

```
One fish, Two fish,
Red Fish, Blue fish.
```

then the output should be:

```
Blue 1
One 1
Red 1
Two 1
Fish, 1
fish, 2
fish. 1
```

As this above output indicates, we did not make any attempt to trim punctuation characters. If we were to make sure that we stripped punctuation and used lowercase characters for each word encountered, we would get this:

```
blue 1
one 1
red 1
two 1
fish 4
```

What follows is a plan for the mapper and reducer functions. You should compare and note the similarity between these and a sequential function for completing this same task on a single input file.

3.2.1 Map-reduce plan

In WMR, mapper functions work simultaneously on lines of input from files, where a line ends with a newline charater. The mapper will produce one key-value pair (*w*, *count*) foreach word encountered in the input line that it is working on.

Thus, on the above input, two mappers working together on each line, after removing punctuation from the end of words and converting the words to lowercase would emit the following combined output:

```
one 1
fish 2
two 1

red 1
fish 2
blue 1
```

The reducers will compute the sum of all the *count* values for a given word *w*, then produce the key-value pair (*w*, *sum*).

3.3 The mapper function

The pseudocode for the mapper looks like this:

```
# key is a single line from a file.
# value is empty in this case, since this is the first mapper function
# we are applying.
#
function mapper(key, value)
  1) Create a dictionary or hash table called counts
  (the keys in counts will be words found and the values will be counts of each word)

  2) Take the key argument to this function, which is the line of text,
     and split it on whitespace

  3) For every word resulting from the split key line:
      strip puntuation from the word
      convert the word to lowercase
      if the word is already a key in the counts dictionary, then
        increment the value in the counts dictioanry by one
      else
        add the key, value pair of (word, 1) to the counts dictionary

  4) For every word 'key' now in the dictionary
      'emit' the (word, count) to the system for the reducers to handle
```

Here is a Python3 mapper function for accomplishing this task using the WMR system. We also add the feature of stripping away puncuation characters from the input.

```python
import string

def mapper(key, value):
    counts = dict()
    words=key.split()
    for word in words:
        word = word.strip(string.punctuation)
        word = word.lower()
        if word not in counts:
            counts[word] = 1
        else:
            counts[word] += 1

    for foundword in counts:
        Wmr.emit(foundword, counts[foundword])
```

This code is available `for download as wc_comb_mapper.py`. You can use this file later when you wish to use it as your mapper in WMR.

Let's examine this code carefully. In line 1 we import the Python `string` package so that we can use its method for returning punctuation characters, found in line 7. Line 3 shows how all mapper functions in WMR should be defined, with two parameters called *key* and *value*. Each of these parameters is a *String* data type. In the case of our first mapper functions reading each line of the file, the whole line is passed into the key in the map-reduce system underlying WMR, and the value is empty. (See additional notes section below for more details you will need when trying other examples.)

In line 4, we create a Python dictionary called *counts* to hold each distinct word and the number of time it appears. In the small input example we describe here, this will not have many entries. When we next read files where a whole book may be contained in one line of data, the dictionary called counts will contain many words.

Line 5 is where we take the input line, which was in the *key*, and break it into words. Then the loop in lines 6-11 goes word by word and strips punctuation and increments the count of that word.

The loop in lines 13 and 14 is how we send the data off to the reducers. The WMR system for Python3 defines a class `Wmr` that includes a class method `emit()` for producing key-value pairs to be forwarded (via shuffling) to a reducer. `Wmr.emit()` requires two string arguments, so both *foundword* and *counts[foundword]* are being interpreted as strings in line 14.

3.4 The reducer function

Pseudocode for a reducer for this problem is exactly the same as for the previous section and looks like this:

```
# key is a single word, values is a 'container' of counts that were
# gathered by the system from every mapper
#
function reducer(key, values)

  set a sum accumulator to zero

  for each count in values
    accumulate the sum by adding count to it

  'emit' the (key, sum) pair
```

A reducer function for solving the word-count problem in Python is

```
1  def reducer(key, values):
2      sum = 0
3      for count in values:
4          sum += int(count)
5      Wmr.emit(key, sum)
```

This code is available `for download as wcreducer.py`. You can use this file later when you wish to use it as your reducer in WMR.

The function `reducer()` is called once for each distinct key that appears among the key-value pairs emitted by the mapper, and that call processes all of the key-value pairs that use that key. On line 1, the two parameters that are arguments of `reducer()` are that one distinct `key` and a Python3 *iterator* (similar to a list, but not quite) called `values`, which provides access to all the values for that key. Iterators in Python3 are designed for `for` loops- note in line 3 that we can simply ask for each value one at a time from the set of values held by the iterator.

Rationale: WMR `reducer()` functions use iterators instead of lists because the number of values may be very large in the case of large data. For example, there would be billions of occurrences of the word "the" if our data consisted of all pages on the web. When there are a lot of key-value pairs, it is more efficient to dispense pairs one at a time through an iterator than to create a gigantic complete list and hold that list in main memory; also, an enormous list may overfill main memory.

The `reducer()` function adds up all the counts that appear in key-value pairs for the `key` that appears as `reducer()`'s first argument (recall these come from separate mappers). Each count provided by the iterator `values` is a string, so in line 4 we must first convert it to an integer before adding it to the running total `sum`.

The method `Wmr.emit()` is used to produce key-value pairs as output from the mapper. This time, only one pair is emitted, consisting of the word being counted and `sum`, which holds the number of times that word appeared in *all* of the original data.

3.5 Running the example code on WebMapReduce

If you have not registered a WMR account or tried the example in the previous section, do that first so that you are used to setting up a job in WMR. You can runn the above mapper and reducer code files on the simple example above in 'Test' mode on WMR to ensure that they work.

3.6 Next Steps

1. In WMR, you can choose to use your own input data files. Try choosing to 'browse' for a file, and using `mobydick.txt` as file input.

2. There are a large number of files of books from Project Gutenberg available on the Hadoop system underlying WebMapReduce. First start by trying this book as an input file by choosing 'Cluster Path' as Input in WMR and typing one of these into the text box:

   /shared/gutenberg/WarAndPeace.txt
   /shared/gutenberg/CompleteShakespeare.txt
   /shared/gutenberg/AnnaKarenina.txt

   These books have many lines of text with 'newline' chacters at the end of each line. Each of you mapper functions works on one line. Try one of these.

3. Next, you should try a collection of many books, each of which has no newline characters in them. In this case, each mapper 'task' in WMR;s underlying Hadoop system will work on one whole book (your dictionary of words per mapper will be for the whole book, and the mappers will be running on many books at one time). *This new version should run faster in principle on the system by letting mappers do a bit of work and pass lass data to the awaiting reducers.* You might want to see if you can see a 'wall clock' time difference bewteen this version and the code described in the previous section. Keep in mind, however, that the time to run depends on how many other users are also using the system.

In the Hadoop file system inside WMR we have these datasets available for this:

'Cluster path' to enter in WMR	Number of books
/shared/gutenberg/all_nonl/group10	2018
/shared/gutenberg/all_nonl/group11	294
/shared/gutenberg/all_nonl/group6	830
/shared/gutenberg/all_nonl/group8	541

While using many books, it will be useful for you to experiment with the different datasets so that you can get a sense for how much a system like Hadoop can process.

To do this, it will also be useful for you to save your configuration so that you can use it again with a different number of reducer tasks. Once you have entered your mapper and reducer code, picked the Python3 language, and given your job a descriptive name, choose the *'Save'* button at the bottom of the WMR panel. This will now be a *'Saved Configuration'* that you can retrieve using the link on the left in the WMR page.

Try using the smallest set first (group11). Do not enter anything in the map tasks box- notice that the system will choose the same number of mappers as the number of books (this will show up once you submit the job). Also do not enter anything for the number of reduce tasks. With that many books, when the job completes you will see there are many pages of output, and some interesting 'words'. For the 294 books in group11, note how you obtain several pages of results. You will also notice that the stripping of punctuation isn't perfect. If you wish to try improving this you could, but it is not necessary.

3.7 Additional Notes

These notes are repeated from the previous section.

It is possible that input data files to mappers may be treated differently than as described in the above example. For example, a mapper function might be used as a second pass by operating on the reducer results from a previous map-reduce cycle. Or the data may be formatted differently than a text file from a novel or poem. These notes pertain to those cases.

In WMR, each line of data is treated as a key-value pair of strings, where the key consists of all characters before the first `TAB` character in that line, and the value consists of all characters after that first `TAB` character. Each call of `mapper()` operates on one line and that function's two arguments are the key and value from that line.

If there are multiple `TAB` characters in a line, all `TAB` characters after the first remain as part of the `value` argument of `mapper()`.

If there are *no* `TAB` characters in a data line (as is the case with all of our fish lines above), an empty string is created for the value and the entire line's content is treated as the key. This is why the key is split in the mapper shown above.

WEBMAPREDUCE ACTIVITIES

This document contains a series of activities for you to try with WebMapReduce. Each one involves separate sets of data of increasing size.

4.1 Poker Hands

The first data set we will explore is about Poker Hands. The Poker Hand Dataset contains a listing of 1,000,000 randomly generated, 5 card poker hands. Each line of the document contains a comma-separated list of the information for each hand. If you read it in order, each line describes first the suit, then the value of each card. The final value on the line is the ranking of the hand (described below).

So, each line representing a particular hand reads, in abbreviated phrasing:

```
'S1,C1,S2,C2,S3,C3,S4,C4,S5,C5,R'
```

where S1 is the value of the suit of card 1 in the hand, C1 is the value of card 1, S2 is the value of the suit of card 2, C2 the value of card 2, and so on.

The suit data values used are:

```
1       Hearts
2       Spades
3       Diamonds
4       Clubs
```

The card data values are:

```
1       Ace
2       Two
3       Three
...
10      Ten
11      Jack
12      Queen
13      King
```

The rankings are numbered in standard poker hand rank ordering:

```
0       High Card
1       Pair
2       Two-Pair
3       Three of a kind
4       Straight (five cards of sequential rank with no gaps)
5       Flush (five cards of the same suit)
```

```
6     Full house (pair + three of a kind)
7     Four of a kind
8     Straight flush (straight of cards all in the same suit)
9     Royal flush (Ace, King, Queen, Jack, Ten, all of the same suit)
```

So, using the mapping of data values on the left to actual cards and hands on the right in the above three tables, a data line for a royal flush in hearts would look like:

```
1,1,1,13,1,12,1,11,1,10,9
```

The first thing we'll do with this is count how many of the 1,000,000 hands are of each ranking. This will involve creating a mapper that splits the inputted string into a list, and then emits the ranking. Then the reducer will count how many times each ranking is emitted. Remember, since this is a comma-separated list and not a tab-separated list, each mapper will get the entire line sent to it as the key, and the value sent to the mapper will be blank. The data has already been uploaded and can be accessed from **/shared/pokerHandData**. Let's start with the mapper, which should look like this:

```python
def mapper(key, value):
    hand = key.split(',')
    rank = hand[10]
    Wmr.emit(rank, 1)
```

We emit a value of 1 to make counting easier in the reducer.

Next, we need our reducer. Remember that since we emitted the rank as the key in the the mapper, each reducer will get a key equal to a rank (0-9), and then an iterator of all of the values emitted by the mappers. Since all of the values emitted by the mappers were 1's, we can simply add up the 1's to get our total count:

```python
def reducer(key, values):
    count = 0
    for value in values:
        count +=int(value)
    Wmr.emit(key, count)
```

4.1.1 Activities

1. Edit the code provided above so that instead of outputting a count, you output the percent of hands in the data set of each ranking.

2. Count the number of flushes of each suit. For a challenge, after you've counted them, convert the suits from their labels to their actual names (change 1 to Hearts, 2 to Spades, and so on).

4.2 Car Information

Next, we'll look at information on cars in another set of comma-separated lists. Each car has six attributes: buying price, maintenance price, number of doors, number of people who fit inside, trunk size, safety, and acceptability of the car. The possible values for each are:

```
Buying Price        v-high, high, med, low
Maintenance Price   v-high, high, med, low
Number of Doors     2, 3, 4, 5-more
Number of People    2, 4, more
Trunk Size          small, med, big
Safety              low, med, high
Acceptability       unacc, acc, good, v-good
```

The data is uploaded at **/shared/carData**

4.2.1 Activities

1. First, using the mapper/reducer you used for the poker hand data, count the number of cars in the set of each acceptability.

2. Adapt your code so that you find a percent of the 1728 cars in the set with a given acceptability.

3. For a challenge, see if you can get more specific, and find cars of a certain acceptability in addition to given attributes. For example, count cars by their price to buy and acceptability. Your output should look something like:

```
unacc-low      Some value
unacc-med      Some value
unacc-high     Some value
acc-low        Some value
acc-med        Some value
        and so on
```

4.3 Movie Data

Next, we're going to look at movie rating data. The information on movie ratings was gathered by a University of Minnesota research group called Movie Lens. The data set contains information on 10,000,054 different ratings, including 10,681 different movies and 71,567 different users (uploaded to **/shared/MovieLens2**). Unlike the previous two datasets, this dataset is arranged into tab-separated lists. Each line contains:

```
MovieId    UserId    Rating    Date
```

For eample, one line of the data looks like this:

```
71567 2338  2 1998-12-01
```

Note: The ratings range in values from 1 to five in 0.5 increments, so consider the ratings to be of data type 'float'.

Before we start playing with the data, let's recall the differences between using a tab-separated list and a comma-separated list. The most obvious difference is using a different split. Instead of splitting on $'$,$'$, we now need to split on $'\backslash t'$. The less obvious difference is how WMR treats the lists. When using a tab-separated list, rather than giving the whole line as the key to the mapper, it gives the first value in the list as the key, and the rest as a single string for the value. In the case of the movie ratings, this means that the key of each mapper will be the MovieId. If it makes it easier for you, you can change the def line of your mapper to read $def\ mapper(movieId,\ value)$.

To make this more clear, let's look at a simple example. Let's count the total number of ratings each movie got. Examine the code below:

```python
def mapper(movieId, value):
    Wmr.emit(movieId, 1)


def reducer(movieId, values):
    count = 0
    for value in values:
        count+= int(value)
    Wmr.emit(movieId, count)
```

4.3.1 Activities

1. Find the average rating for each movie.

2. Find the average rating that each user gives to movies.

3. Find the number of movies given each of the five ratings.

4. Find the average rating per year

5. Find the average rating in July of each year

4.4 Flight Data

Provided by the Bureau of Transportation Statistics, the Flight Data dataset (the data is uploaded to **/shared/FlightData**) contains information on delayed and cancelled flights. Each line of the data is arranged in a comma-separated list detailing: Flight Date, Airline, Origin Airport, Origin State, Destination Airport, Destination State, Departure Delay, Arrival Delay, Cancellation Code, Carrier Delay, Weather Delay, Security Delay, Late Aircraft Delay, Totally Additional Gate Time.

A couple notes about the data. First, notice that a negative delay means an early departure or arrival. Also, it is important to note that all of the text entries in the data include quotes. Numbers are represented in floating point without quotes. If you want to include quotes in a string, you need to use a backlash. You can also use the strip() method on any string to remove leading and trailing characters. So if you $import\ string$, you can do $strip(string.punctuation)$ to remove all punctuation, including quotation marks, from the string. Next, be careful with cancelled flights. Cancelled flights are represented differently in the data than flights that were simply delayed, in that the delay is left blank, but a code is put in the Cancellation Code column. This means that somewhere in your mapper or reducer you have to have a condition to deal with these, or else your values will not come out well.

The data is organized into 4 folders. Each folder represents a year's worth of information. Thus, within **/shared/FlightData**, are directories for data from 2011, 2010, 2009, 2008. Each file in those folders contains a month's worth of information. Each one of these has files for each month of that year. So to get the January 2011 data, your Cluster Path would be:

/shared/FlightData/2011/201101.csv.

Note: that if you type **/shared/FlightData/2011** into the Cluster Data Path, you will use all of the files for the year 2011 (the 2001 data is an incomplete set, in that it contains data from January through April). Thus, you can do a year's worth of data at a time.

Note: A nice trick when using WebMapReduce is that you can choose the test option on one month's worth of data and enter the identity mapper and reducer to simply get a sense for what is in the first few lines of the file itself. (Ask your instructor if you do not have example Python files for an identity mapper and an identity reducer.)

Do this now: use an identity mapper and identity reducer on this file:

/shared/FlightData/2011/201101.csv

Note how the date is formatted: the date string is "year-month-day" as "yyyy-mm-dd". So a flight on January 1, 2011 has a date string "2011-01-01".

A potential issue: Now that we've mentioned the nice trick about using test mode, it is sometimes tha case that test mode seems to stop working in WMR. When this happens, you are left to simply submit your work instead.

There is a 'header' line in each file that indicates what is in each 'column' of data separated by the commas. It looks like this (all on one line in the file):

```
"FL_DATE","CARRIER","ORIGIN","ORIGIN_STATE_ABR","DEST","DEST_STATE_ABR","DEP_DELAY","
ARR_DELAY","CANCELLATION_CODE","CARRIER_DELAY","WEATHER_DELAY","NAS_DELAY",
"SECURITY_DELAY","LATE_AIRCRAFT_DELAY","TOTAL_ADD_GTIME",
```

To see just this line, you could use a mapper like this:

```python
def mapper(key, value):
    items = key.split(',')
    if items[0] =='"FL_DATE"':
        Wmr.emit(key, value)
```

Then use an 'identity reducer' with the above.

Now you have seen what is in this file. Before you can use this data for analysis, you must first add a condition into your mapper that deals with the first line of the file. If you examine the files, you will see that the first line of each file is a header file that details what information is on each line of the file. This makes the file a lot easier to read, and is especially useful if you are using the python csv module (which we will not use in WebMapReduce). In our case however, you need to put a condition in your mapper to ignore this line. Think about this: if you split the key, what will the first element in the list be? Will it ever be the same thing in any of the other lines as it is in the first line?

4.4.1 Activities

1. First, pick a year and find the average arrival delay for each airport in that year. Use the origin airport.

2. *in homework:* Find the average arrival delay per day.

3. *Challenge:* Find the average arrival delay per month. Hint: While similar to finding the average delay per day, this involves an extra step.

4. *Challenge:* find the average delay per airline per month. To do this, you will have to run jobs for one airport at a time. Pick specific airlines to try. Start with the major ones like Delta (DL), United (UA), American (AA), or Southwest (WN). A note about using this data with Google Fusion Tables: to get Fusion Tables to recognize a month as a month, you need to have it in the form $mm/yyyy$. This means you have to split the date string as it is given, and then create a new string using a slash (/) instead of a dash (-).

4.5 Google N-Grams

A N-Gram is a phrase of n words. For example, "hello" is a 1-Gram, and "hello world" is a 2-Gram, or bi-gram. Using books from Google Books, Google put together a list of N-Grams. Last generated in July 2009, the corpora contains 10 Gigabytes(GB) of 1-Grams, 100 GB of 2-Grams, and 200 GB of 3-Grams. 4-Grams and 5-Grams are also available. The n-Grams data is uploaded to **/shared/NGrams**

The N-Grams are arranged into files which contain tab-separated lists. Each line shows the information for an N-Gram for a given year. It gives the following information:

```
N-Gram   Year   Total occurrences   Pages   Volumes
```

The $Pages$ entry is the total number of pages an N-Gram occurs on. So if the word **and** appears 5 times on a page in a book, it counts 5 times for the $Total\ occurences$, but only once for the $Pages$. $Volumes$ is the same thing as $Pages$, except that it counts the number of unique volumes or books that each N-Gram occurs in.

4.5.1 1-Grams

Let's start by working with a useful 1-Grams activity. An interesting problem you can investigate with N-Grams is how language has developed over time. As language evolves, new words enter peoples' vocabulary, while others fall

into obscurity. I'm sure you can think of many examples, like how *thou* has fallen into obscurity, while the word *computer* is a relatively modern word.

We will look at a useful method for determining high-frequency interesting words. Our goal is to eliminate highly occurring words of low interest, such as articles (the, a) and prepositions (e.g. to, from, of for) and focus on 'interesting' words that occur often.

Information retrieval experts are interested in a related problem: given a set of documents and a user's query word, find all those related to that particular word. This is done by locating the documents where that word occurs the most often in relation to the size of the document and number of total documents.

We can use this technique in a slightly different way to determine the frequency of popular words, yet eliminate those that are simply commonly occurring words in English. We will do this with a ratio called a *tf-idf* (term frequency-inverse document frequency). The formula for tf-idf is:

$$log\left(\frac{number\ of\ documents\ that\ year}{number\ of\ documents\ the\ word\ appears\ in}\right) * Total\ occurrences$$

Notice how the fraction approaches one for uninteresting words that occur in every document. Since the log of one is zero, this value will be quite low. Those words that occur more frequently, but not in every document, will have higher values. We will be examining how to use this to determine some of the top frequently occurring words per year in the 1-grams dataset.

There are some Python files to help you get started in a directory of files on moodle.

The file called *1gramMap.py* in this directory contains a dictionary called *yearDict* that has defined the total number of unique 1-Grams for each year. We did some separate analysis of the 1-grams to devise this dictionary for you. What is this mapper emitting? Note that we are eliminating years where there is not very much data (low number of volumes), because the tf-idf calculation is less useful for these.

Now let us examine the reducer, in a file called *1gramReduce.py*. Look it over and explain what it is doing. Write explanations as comments in each of these files.

There are likely a few new things in this code that you have not seen before. OIne of them is the use of the *sorted* method to sort the items in a dictionary. Try to look up how this works. We need to sort the words out into a list of pairs (word, frequency), ordered by frequency, in order to emit only the top 20 frequently occurring words. If you still find this confusing, try creating a simple example and using it in a script. Create a dictionary that is not sorted, sort it, and then loop through it getting the sorted values.

Activity

Run this code through the 1-gram data, found on **/shared/NGrams/1-Grams**. Do you find anything interesting in any particular year?

Activity

Now you will do something different. Your goal: pick some words and see how their frequencies have changed over time. You can do this one word at a time. Your final output will be the pairs of (year, frequency of that word in that year).

In your mapper, you will want to ignore all words in the files, except for the word you are looking for. By doing this, we are using the many mappers that will each work on one line of this data as 'filters' to eliminate all words but the one we are interested in. In this case the year will be the key. You decide

Each reducer will work on a year. In your reducer, you can use a dictionary of the total number of words per year in the collection. We have created this for you in the file `findWordReduce.py`. Figure out how this will help you generate the frequency of the word in that year.

Do These

1. Choose some words that were given as examples in the 'Culturomics' paper.

2. Try graphing your results for a visual representation (see Google Fusion tables below).

3. You could compare words to see how related they are (e.g Microsoft, Apple or computer, technology).

4.6 Google Fusion Tables

In this section, we will explore using the Google Labs project Google Fusion Tables. Google Fusion Tables allows you to import tables of data, merge them together, and then create visualizations. If you follow the below steps, you should be able to get started very easily.

1. First, you need a Google account. It appears that this will not work with your Macalester Google Apps account. You will need to use it with another gmail address ending in @gmail.com. If you don't have one, you could create a Google account. They are free and easy to set up.

2. Go to the bar at the top of any google page which lists the different Google Services. The last item should read *more*. Click on it, then click on *even more* at the bottom of the dropdown menu.

3. Scroll to the bottom of the page. Under the last heading, $Innovation$, you should see a link to Google Fusion Tables. Click on it.

4. From the main page, you can see a list of all the tables you have created/imported (using a gmail address).

    If you see a list of all of your files from your Macalester Google Apps account, this means you are logged into that, and will need to log out. Or you can use a different browser and log into your gmail account with that.

5. Whe logged into your Gmail account: You can view your Google Docs available with that account by choosing 'View My Tables'. Here is Google's help page for importing your files for use with Fusion Tables:

    Google Fusion Table Help to Import Files

    You will be importing your files into Google Docs as 'Tables'. Be sure to choose that your file is tab-separated if it came from WMR.

    (See example files below if you don't have one of your own too try.)

6. To create a graph of your data table now in Google Docs, simply select $Visualize$ from the top menu of your table. Then, select the type of visualization you would like to use.

7. Play around with the different features of the visualizations to get a feel of how to create graphs.

4.7 Merging Tables

There are two files that you can try using for this available on moodle:

1. ComputerOccurences.csv

2. technologyOccurences.csv

Though the names imply that they are comma-separated, the data in each column is actually tab-separated.

Some of the activities you will do will require you to merge output from a couple of WMR job outputs like these together and visualize them together. Doing this is quite simple:

1. Open one of the files you would like to merge from the Google Fusion Talbes main page.

2. Select $Merge$ from the top menu of the table.

3. Click in the input box next to the big number 2. This should open a drop-down list of your other tables. Select the one you would like to merge to your table.

4. In the two side-by-side boxes below this, click on the radio-button (it looks like a grey dot) next to the value that is shared between both tables. This should be something like year, or month, or whatever the common variable is between the two tables.

5. Below both boxes, check the $Select\ columns$ option. Then, make sure that all of the square boxes are checked next to your variables.

6. Enter a name for the new table you are making in the text box where it asks you to.

7. Click $Merge\ tables$

8. You can then visualize the new table just like any other table. To select to show all values, either hold the CTRL or SHIFT button on your keyboard as you select the columns in the list.