

---

# Parallel Patternlets

**CSinParallel Project**

July 26, 2013

# CONTENTS

|          |  |          |
|----------|--|----------|
| <b>1</b> | <b>Source Code</b>                                     | <b>2</b> |
| <b>2</b> | <b>Patternlet Examples</b>                             | <b>3</b> |
| 2.1      | Message Passing Parallel Patternlets . . . . .         | 3        |
| 2.2      | Shared Memory Parallel Patternlets in OpenMP . . . . . | 16       |

This document contains simple examples of basic elements that are combined to form patterns often used in programs employing parallelism. The examples are separated between two major *coordination patterns*:

1. message passing used on single multiprocessor machines or clusters of distributed computers, and
2. mutual exclusion between threads executing concurrently on a single shared memory system.

Both sets of examples are illustrated with the C programming language, using standard popular available libraries. The message passing example uses a C library called MPI (Message Passing Interface). The mutual Exclusion/shared memory examples use the OpenMP library.

# SOURCE CODE

Please download all examples from this tarball: `patternlets.tgz`

# PATTERNLET EXAMPLES

## 2.1 Message Passing Parallel Patternlets

Parallel programs contain *patterns*: code that recurs over and over again in solutions to many problems. The following examples show very simple examples of small portions of these patterns that can be combined to solve a problem. These C code examples use the Message Passing Interface (MPI) library, which is suitable for use on either a single multiprocessor machine or a cluster of machines.

### 2.1.1 Source Code

Please download all examples from this tarball: `patternlets.tar.gz`

A C code file for each example below can be found in subdirectories of the MPI directory, along with a makefile and an example of how to execute the program.

### 2.1.2 0. Single Program, Multiple Data

First let us illustrate the basic components of an MPI program, which by its nature uses a single program that runs on each process. Note what gets printed is different for each process, thus the processes using this one single program can have different data values for its variables. This is why we call it single program, multiple data.

```
/* spmd.c
 * ... illustrates the single program multiple data
 *      (SPMD) pattern using basic MPI commands.
 *
 * Joel Adams, Calvin College, November 2009.
 *
 * Usage: mpirun -np 4 ./spmd
 */

#include <stdio.h>
#include <mpi.h>

int main(int argc, char** argv) {
    int id = -1, numProcesses = -1, length = -1;
    char myHostName[MPI_MAX_PROCESSOR_NAME];

    MPI_Init(&argc, &argv);
    MPI_Comm_rank(MPI_COMM_WORLD, &id);
    MPI_Comm_size(MPI_COMM_WORLD, &numProcesses);
    MPI_Get_processor_name(myHostName, &length);
```

```
    printf("Greetings from process %d of %d on %s\n",
           id, numProcesses, myHostName);

    MPI_Finalize();
    return 0;
}
```

file: patternlets/MPI/00.spmc/spmc.c

### 2.1.3 1. The Master-Worker Implementation Strategy Pattern

```
/* masterServer.c
 * ... illustrates the basic master-worker pattern in MPI ...
 * Joel Adams, Calvin College, November 2009.
 */

#include <stdio.h>
#include <mpi.h>

int main(int argc, char** argv) {
    int id = -1, numWorkers = -1, length = -1;
    char hostName[MPI_MAX_PROCESSOR_NAME];

    MPI_Init(&argc, &argv);
    MPI_Comm_rank(MPI_COMM_WORLD, &id);
    MPI_Comm_size(MPI_COMM_WORLD, &numWorkers);
    MPI_Get_processor_name (hostName, &length);

    if ( id == 0 ) { // process 0 is the master
        printf("Greetings from the master, # %d (%s) of %d processes\n",
               id, hostName, numWorkers);
    } else { // processes with ids > 0 are workers
        printf("Greetings from a worker, # %d (%s) of %d processes\n",
               id, hostName, numWorkers);
    }

    MPI_Finalize();
    return 0;
}
```

file: patternlets/MPI/01.masterWorker/masterWorker.c

### 2.1.4 2. Message passing 1, using Send-Receive of a single value

```
/* messagePassing.c
 * ... illustrates the use of the MPI_Send() and MPI_Recv() commands...
 * Joel Adams, Calvin College, November 2009.
 */

#include <stdio.h>
#include <mpi.h>
#include <math.h> // sqrt()

int odd(int number) { return number % 2; }
```

```

int main(int argc, char** argv) {
    int id = -1, numProcesses = -1;
    float sendValue = -1, receivedValue = -1;
    MPI_Status status;

    MPI_Init(&argc, &argv);
    MPI_Comm_rank(MPI_COMM_WORLD, &id);
    MPI_Comm_size(MPI_COMM_WORLD, &numProcesses);

    if (numProcesses > 1 && !odd(numProcesses) ) {
        sendValue = sqrt(id);
        if ( odd(id) ) { // odd processors send, then receive
            MPI_Send(&sendValue, 1, MPI_FLOAT, id-1, 1, MPI_COMM_WORLD);
            MPI_Recv(&receivedValue, 1, MPI_FLOAT, id-1, 2,
                    MPI_COMM_WORLD, &status);
        } else { // even processors receive, then send
            MPI_Recv(&receivedValue, 1, MPI_FLOAT, id+1, 1,
                    MPI_COMM_WORLD, &status);
            MPI_Send(&sendValue, 1, MPI_FLOAT, id+1, 2, MPI_COMM_WORLD);
        }

        printf("Process %d of %d computed %f and received %f\n",
               id, numProcesses, sendValue, receivedValue);
    } else if ( !id ) { // only process 0 does this part
        printf("\nPlease run this program using -np N where N is positive and even.\n\n");
    }

    MPI_Finalize();
    return 0;
}

```

file: patternlets/MPI/02.messagePassing/messagePassing.c

### 2.1.5 3. Message passing 2, using Send-Receive of an array of values

```

/* messagePassing2.c
 * ... illustrates the use of the MPI_Send() and MPI_Recv() commands...
 *
 * Joel Adams, Calvin College, November 2009.
 *
 * Usage: mpirun -np N ./messagePassing2
 *
 * Exercise: Run the program, varying the value of N from 1-8.
 */

#include <stdio.h> // printf()
#include <string.h> // strlen()
#include <mpi.h> // MPI

#define MAX 256

int main(int argc, char** argv) {
    int id = -1, numProcesses = -1;
    char sendBuffer[MAX] = {'\0'};
    char recvBuffer[MAX] = {'\0'};
    MPI_Status status;

```

```

MPI_Init(&argc, &argv);
MPI_Comm_rank(MPI_COMM_WORLD, &id);
MPI_Comm_size(MPI_COMM_WORLD, &numProcesses);

if (numProcesses > 1) {
    if ( id == 0 ) {
        sprintf(sendBuffer, "%d", id);    // create msg

        MPI_Send(sendBuffer,                // msg sent
                  strlen(sendBuffer) + 1,  // num chars + NULL
                  MPI_CHAR,                 // type
                  id+1,                     // destination
                  1,                        // tag
                  MPI_COMM_WORLD);          // communicator

        MPI_Recv(recvBuffer,                // msg received
                  MAX,                      // buffer size
                  MPI_CHAR,                 // type
                  numProcesses-1,           // sender
                  1,                        // tag
                  MPI_COMM_WORLD,           // communicator
                  &status);                 // recv status
    } else {
        MPI_Recv(recvBuffer,                // msg received
                  MAX,                      // buffer size
                  MPI_CHAR,                 // type
                  MPI_ANY_SOURCE,           // sender (anyone)
                  1,                        // tag
                  MPI_COMM_WORLD,           // communicator
                  &status);                 // recv status

        // build msg to send by appending id to msg received
        sprintf(sendBuffer, "%s %d", recvBuffer, id);

        MPI_Send(sendBuffer,                // msg to send
                  strlen(sendBuffer) + 1,  // num chars + NULL
                  MPI_CHAR,                 // type
                  (id+1) % numProcesses,    // destination
                  1,                        // tag
                  MPI_COMM_WORLD);          // communicator
    }

    printf("Process %d of %d received %s\n",
           id, numProcesses, recvBuffer);
} else {
    printf("\nPlease run this program with at least 2 processes\n\n");
}

MPI_Finalize();
return 0;
}

```

file: patternlets/MPI/03.messagePassing2/messagePassing2.c



### 2.1.6 4. A. Data Decomposition: on *slices* using parallel-for (textual version)

In this example, the data being decomposed is simply the set of integers from zero to  $\text{REPS} * \text{numProcesses}$ , which are used in the for loop.

```
/* parallelForSlices.c
 * ... illustrates the parallel for loop pattern in MPI
 *      in which processes perform the loop's iterations in 'slices'
 *      (simple, and useful when loop iterations do not access
 *      memory/cache locations) ...
 * Joel Adams, Calvin College, November 2009.
 */

#include <stdio.h>
#include <mpi.h>

int main(int argc, char** argv) {
    const int REPS = 8;
    int id = -1, numProcesses = -1, i = -1;

    MPI_Init(&argc, &argv);
    MPI_Comm_rank(MPI_COMM_WORLD, &id);
    MPI_Comm_size(MPI_COMM_WORLD, &numProcesses);

    for (i = id; i < REPS; i += numProcesses) {
        printf("Process %d is performing iteration %d\n",
              id, i);
    }

    MPI_Finalize();
    return 0;
}
```

file: patternlets/MPI/04.parallelForLoop-slices/textual/parallelForSlices.c

### 2.1.7 4. B. Data Decomposition: on *slices* using parallel-for (visual version)

In this example, we can visually see how the slicing of data used in iterations of a nested for loop is working. Run it to see the effect!

```
/* parallelForStripes.c is a graphical illustration of
 * the 'slicing' version of the parallel for loop design pattern,
 * using Argonne Labs' MPE graphics library for X11 systems.
 *
 * Summer 2013, Joel Adams, Calvin College.
 *
 * Usage: mpirun -np N ./parallelForStripes
 *      Click the mouse in the window to terminate the program.
 *      You must have an X11 server running.
 *
 * Exercise: Run the program, varying N from 1 - 32, like this: 1, 2, 4, 8, 16, 24, 32
 *      and compare to the 'Blocks' version...
 */

#include <mpi.h>      // MPI
#include <mpe.h>      // MPE
#include <stdlib.h>    // getenv()
```

---

```

#include <string.h> // strcmp()
#include <stdio.h> // printf(), etc.
#include <unistd.h> // usleep

/*
 * getDisplay() retrieves the DISPLAY environment info
 */

char* getDisplay() {
    char * display = getenv("DISPLAY");
    if ( strcmp(display, "(null)", 7) == 0 ) {
        fprintf(stderr, "\n*** Fatal: DISPLAY variable not set.\n");
        exit(1);
    }
    return display;
}

int main(int argc, char* argv[]) {
    const int WINDOW_WIDTH = 800;
    const int WINDOW_HEIGHT = 512;
    int x = 0, y = 0;
    int id = -1, numProcesses = -1;
    int button = -1;
    int square_width = 4;
    MPE_XGraph canvas;
    MPE_Color *colors = NULL;
    MPE_Color myColor = 0;

    // initialize environment, variables, etc.
    MPI_Init(&argc,&argv);
    MPI_Comm_rank(MPI_COMM_WORLD, &id);
    MPI_Comm_size(MPI_COMM_WORLD, &numProcesses);
    MPE_Open_graphics( &canvas, MPI_COMM_WORLD,
                      getDisplay(),
                      5, 5,
                      WINDOW_WIDTH, WINDOW_HEIGHT, 0 );
    colors = malloc( numProcesses * sizeof(MPE_Color) );
    MPE_Make_color_array(canvas, numProcesses, colors);
    myColor = colors[id];

    // each process does a block of the window's pixels
    ///////////////////////////////////////////////////
    // You could try it using the following commented code and more processes
    // This version draws single points for each execution in the nested loop
    //for (y = id; y < WINDOW_HEIGHT; y += numProcesses) {
    //    for (x = 0; x < WINDOW_WIDTH; x++) {
    //        MPE_Draw_point(canvas, x, y, myColor);
    //        usleep(4000); // for illustrative purposes, we wait so our
    //                      // eyes can see it happening
    //        MPE_Update(canvas); // update here to slow things down
    //    }
    //}
    ///////////////////////////////////////////////////
    // This version draws bars 4 pixels wide and waits so we can see it
    for (y = id*square_width; y < WINDOW_HEIGHT; y += numProcesses*square_width) {
        for (x = 0; x < WINDOW_WIDTH; x+=square_width) {
            MPE_Fill_rectangle( canvas, x, y, square_width, square_width, myColor );
            usleep(4000); // for illustrative purposes,
                          // we wait so our eyes can see it happening
        }
    }
}

```

```

        MPE_Update(canvas); // update here to slow things down
    }
}

// pause until mouse-click so the program doesn't terminate
if (id == 0) {
    printf("\nClick in the window to continue...\n");
    MPE_Get_mouse_press( canvas, &x, &y, &button );
}

// clean up
MPE_Close_graphics( &canvas);
MPI_Finalize();
free(colors);
if (id == 0) printf("Program complete.\n\n");
return 0;
}

```

file: *patternlets/MPI/04.parallelForLoop-slices/visual/parallelForSlices.c*

## 2.1.8 5. A. Data Decomposition: on *blocks* using parallel-for (textual version)

This is a basic example that does not yet include a data array, though it would typically be used when each process would be working on a portion of an array that could have been looped over in a sequential solution.

```

/* parallelForBlocks.c
 * ... illustrates the parallel for loop pattern in MPI
 *      in which processes perform the loop's iterations in 'blocks'
 *      (preferable when loop iterations do access memory/cache locations) ...
 * Joel Adams, Calvin College, November 2009.
 */

#include <stdio.h>
#include <mpi.h>

int main(int argc, char** argv) {
    const int REPS = 8;
    int id = -1, numProcesses = -1, i = -1,
        start = -1, stop = -1, blockSize = -1;

    MPI_Init(&argc, &argv);
    MPI_Comm_rank(MPI_COMM_WORLD, &id);
    MPI_Comm_size(MPI_COMM_WORLD, &numProcesses);

    blockSize = REPS / numProcesses; // integer division
    start = id * blockSize;           // find starting index
                                     // find stopping index
    if ( id < numProcesses - 1 ) {   // if not the last process
        stop = (id + 1) * blockSize; // stop where next process starts
    } else {                          // else
        stop = REPS;                 // last process does leftovers
    }

    for (i = start; i < stop; i++) { // iterate through your range
        printf("Process %d is performing iteration %d\n",
            id, i);
    }
}

```

```

    MPI_Finalize();
    return 0;
}

```

file: *patternlets/MPI/05.parallelForLoop-blocks/textual/parallelForBlocks.c*

### 2.1.9 5. B. Data Decomposition: on *blocks* using parallel-for (visual version)

In this example you can see how blocks of values within a matrix might be assigned to each process. Run it to see the effect!

```

/* parallelForBlocks.c is a graphical illustration of
 * the 'blocks' version of the parallel for loop design pattern,
 * using Argonne Labs' MPE graphics library for X11 systems.
 *
 * Summer 2013, Joel Adams, Calvin College.
 *
 * Usage: mpirun -np N ./parallelForBlocks
 *        Click the mouse in the window to terminate the program.
 *        You must have an X11 server running.
 *
 * Exercise: Run the program, varying N from 1 - 32, like this: 1, 2, 4, 8, 16, 24, 32
 *           and compare to the 'Stripes' version...
 */

#include <mpi.h>      // MPI
#include <mpe.h>      // MPE
#include <stdlib.h>    // getenv()
#include <string.h>    // strcmp()
#include <stdio.h>     // printf(), etc.
#include <unistd.h>    // usleep

/*
 * getDisplay() retrieves the DISPLAY environment info
 */

char* getDisplay() {
    char * display = getenv("DISPLAY");
    if ( strcmp(display, "(null)", 7) == 0 ) {
        fprintf(stderr, "\n*** Fatal: DISPLAY variable not set.\n");
        exit(1);
    }
    return display;
}

int main(int argc, char* argv[]) {
    const int WINDOW_WIDTH = 800;
    const int WINDOW_HEIGHT = 512;
    int x = 0, y = 0;
    int id = -1, numProcesses = -1;
    int button = -1;
    int blockSize = 0;
    int yStart, yEnd;
    MPE_XGraph canvas;
    MPE_Color *colors = NULL;
    MPE_Color myColor = 0;

```

```

// initialize environment, variables, etc.
MPI_Init(&argc,&argv);
MPI_Comm_rank(MPI_COMM_WORLD, &id);
MPI_Comm_size(MPI_COMM_WORLD, &numProcesses);
MPE_Open_graphics( &canvas, MPI_COMM_WORLD,
                  getDisplay(),
                  5, 5,
                  WINDOW_WIDTH, WINDOW_HEIGHT, 0 );
colors = malloc( numProcesses * sizeof(MPE_Color) );
MPE_Make_color_array(canvas, numProcesses, colors);
myColor = colors[id];
blockSize = WINDOW_HEIGHT / numProcesses;
yStart = id * blockSize;
yEnd = yStart + blockSize;

// processes draw the window's pixels in stripes
// You could try it using the following commented code and more processes
// This version draws single points for each execution in the nested loop
// for (y = yStart; y < yEnd; y++) {
//     for (x = 0; x < WINDOW_WIDTH; x++) {
//         MPE_Draw_point(canvas, x, y, myColor);
//         usleep(4000); // for illustrative purposes, we wait so our
//                        // eyes can see it happening
//         MPE_Update(canvas); // update here to slow things down
//     }
// }
// This version draws bars 4 pixels wide and waits so we can see it
for (y = yStart; y < yEnd; y+=4) {
    for (x = 0; x < WINDOW_WIDTH; x+=4) {
        MPE_Fill_rectangle( canvas, x, y, 4, 4, myColor );
        usleep(4000); // for illustrative purposes, we wait so our
                      // eyes can see it happening
        MPE_Update(canvas); // update here to slow things down
    }
}

// pause until mouse-click so the program doesn't terminate
if (id == 0) {
    printf("\nClick in the window to continue...\n");
    MPE_Get_mouse_press( canvas, &x, &y, &button );
}

// clean up
MPE_Close_graphics( &canvas);
MPI_Finalize();
free(colors);
if (id == 0) printf("Program complete.\n\n");
return 0;
}

```

file: *atternlets/MPI/05.parallelForLoop-blocks/visual/parallelForBlocks.c*

## 2.1.10 6. Broadcast: a special form of message passing

This example shows how to ensure that all processes have a copy of an array created by a single *master* node.

```

/* bcast.c
 * ... illustrates the use of MPI_Bcast()...
 * Joel Adams, Calvin College, November 2009.
 */

#include <mpi.h>
#include <stdio.h>
#include <stdlib.h>

/* fill an array with values
 */
void fill(int* a, int size) {
    int i;
    for (i = 0; i < size; i++) {
        a[i] = i+11;
    }
}

/* display a string, a process id, and its array values
 */
void print(char* str, int id, int* a) {
    printf("process %d %s: {%d, %d, %d, %d, %d, %d, %d, %d}\n",
        id, str, a[0], a[1], a[2], a[3], a[4], a[5], a[6], a[7]);
}

#define MAX 8

int main(int argc, char** argv) {
    int array[MAX] = {0};
    int numProcs, myRank;

    MPI_Init(&argc, &argv);
    MPI_Comm_size(MPI_COMM_WORLD, &numProcs);
    MPI_Comm_rank(MPI_COMM_WORLD, &myRank);

    if (myRank == 0) fill(array, MAX);

    print("array before", myRank, array);
    MPI_Bcast(array, MAX, MPI_INT, 0, MPI_COMM_WORLD);
    print("array after", myRank, array);

    MPI_Finalize();

    return 0;
}

```

*file: patternlets/MPI/06.broadcast/broadcast.c*

## 2.1.11 7. Collective Communication: Reduction

Once processes have performed independent concurrent computations, possibly on some portion of decomposed data, it is quite common to then *reduce* those individual computations into one value. This example shows a simple calculation done by each process being reduced to a sum and a maximum. In this example, MPI, has built-in computations, indicated by MPI\_SUM and MPI\_MAX in the following code.

```

/* reduce.c
 * ... illustrates the use of MPI_Reduce()...

```

```
 * Joel Adams, Calvin College, November 2009.
 */

#include <mpi.h>
#include <stdio.h>
#include <stdlib.h>

int main(int argc, char** argv) {
    int numProcs = -1, myRank = -1, square = -1, max = -1, sum = 0;

    MPI_Init(&argc, &argv);
    MPI_Comm_size(MPI_COMM_WORLD, &numProcs);
    MPI_Comm_rank(MPI_COMM_WORLD, &myRank);

    square = (myRank+1) * (myRank+1);

    printf("Process %d computed %d\n", myRank, square);

    MPI_Reduce(&square, &sum, 1, MPI_INT, MPI_SUM, 0, MPI_COMM_WORLD);

    MPI_Reduce(&square, &max, 1, MPI_INT, MPI_MAX, 0, MPI_COMM_WORLD);

    if (myRank == 0) {
        printf("\nThe sum of the squares is %d\n\n", sum);
        printf("The max of the squares is %d\n\n", max);
    }

    MPI_Finalize();

    return 0;
}
```

*file: patternlets/MPI/07.reduction/reduction.c*

## 2.1.12 8. Collective communication: Scatter for message-passing data decomposition

If processes can independently work on portions of a larger data array using the geometric data decomposition pattern, the scatter patternlet can be used to ensure that each process receives a copy of its portion of the array.

```
/* scatter.c
 * ... illustrates the use of MPI_Scatter()...
 * Joel Adams, Calvin College, November 2009.
 */

#include <mpi.h>
#include <stdio.h>
#include <stdlib.h>

int main(int argc, char** argv) {
    const int MAX = 8;
    int aSend[] = {11, 22, 33, 44, 55, 66, 77, 88};
    int* aRcv;
    int i, numProcs, myRank, numSent;

    MPI_Init(&argc, &argv);
    MPI_Comm_size(MPI_COMM_WORLD, &numProcs);
```

```

MPI_Comm_rank(MPI_COMM_WORLD, &myRank);

numSent = MAX / numProcs;
aRcv = (int*) malloc( numSent * sizeof(int) );
MPI_Scatter(aSend, numSent, MPI_INT, aRcv, numSent, MPI_INT, 0, MPI_COMM_WORLD);
printf("Process %d: ", myRank);
for (i = 0; i < numSent; i++) {
    printf(" %d", aRcv[i]);
}
printf("\n");

free(aRcv);
MPI_Finalize();

return 0;
}

```

*file: patternlets/MPI/08.scatter/scatter.c*

### 2.1.13 9. Collective communication: Gather for message-passing data decomposition

If processes can independently work on portions of a larger data array using the geometric data decomposition pattern, the gather patternlet can be used to ensure that each process sends a copy of its portion of the array back to the root, or master process.

```

/* gather.c
 * ... illustrates the use of MPI_Gather()...
 * Joel Adams, Calvin College, November 2009.
 */

#include <mpi.h>
#include <stdio.h>
#include <stdlib.h>

#define MAX 3

int main(int argc, char** argv) {
    int computeArray[MAX];
    int* gatherArray = NULL;
    int i, numProcs, myRank, totalGatheredVals;

    MPI_Init(&argc, &argv);
    MPI_Comm_size(MPI_COMM_WORLD, &numProcs);
    MPI_Comm_rank(MPI_COMM_WORLD, &myRank);

    for (i = 0; i < MAX; i++) {
        computeArray[i] = myRank * 10 + i;
    }

    totalGatheredVals = MAX * numProcs;
    gatherArray = (int*) malloc( totalGatheredVals * sizeof(int) );

    MPI_Gather(computeArray, MAX, MPI_INT,
              gatherArray, MAX, MPI_INT, 0, MPI_COMM_WORLD);

    if (myRank == 0) {

```



```

    for (i = 0; i < totalGatheredVals; i++) {
        printf(" %d", gatherArray[i]);
    }
    printf("\n");
}

free(gatherArray);
MPI_Finalize();

return 0;
}

```

*file: patternlets/MPI/09.gather/gather.c*

## 2.1.14 10. Collective Communication: Barrier

This simple example shows the use of a barrier: a point at which all processes must complete the code above it before moving on and running the code below it.

```

/* barrier.c
 * ... illustrates the behavior of MPI_Barrier() ...
 * Joel Adams, Calvin College, May 2013.
 *
 * Exercise:
 * - Compile and run the program several times, noting the intermixed outputs
 * - Uncomment the barrier call; then recompile and rerun,
 *   noting the change in the outputs.
 */

#include <stdio.h>
#include <mpi.h>

int main(int argc, char** argv) {
    int id = -1, numProcesses = -1, length = -1;
    char myHostName[MPI_MAX_PROCESSOR_NAME];

    MPI_Init(&argc, &argv);
    MPI_Comm_rank(MPI_COMM_WORLD, &id);
    MPI_Comm_size(MPI_COMM_WORLD, &numProcesses);
    MPI_Get_processor_name(myHostName, &length);

    printf("Process %d of %d on %s is BEFORE the barrier.\n",
           id, numProcesses, myHostName);

    MPI_Barrier(MPI_COMM_WORLD);

    printf("Process %d of %d on %s is AFTER the barrier.\n",
           id, numProcesses, myHostName);

    MPI_Finalize();
    return 0;
}

```

*file: patternlets/MPI/10.barrier/barrier.c*

## 2.2 Shared Memory Parallel Patternlets in OpenMP

When writing programs for shared-memory hardware with multiple cores, a programmer could use a low-level thread package, such as pthreads. An alternative is to use a compiler that processes OpenMP *pragmas*, which are compiler directives that enable the compiler to generate threaded code. Whereas pthreads uses an **explicit** multithreading model in which the programmer must explicitly create and manage threads, OpenMP uses an **implicit** multithreading model in which the library handles thread creation and management, thus making the programmer's task much simpler and less error-prone.

The following are examples of C code with OpenMP pragmas. The first three are basic illustrations to get used to the OpenMP pragmas. The rest illustrate how to implement particular patterns and what can go wrong when mutual exclusion is not properly ensured.

Note: by default OpenMP uses the **Thread Pool** pattern of concurrent execution. OpenMP programs initialize a group of threads to be used by a given program (often called a pool of threads). These threads will execute concurrently during portions of the code specified by the programmer.

### 2.2.1 Source Code

Please download all examples from this tarball: `patternlets.tar.gz`

A C code file for each example below can be found in subdirectories of the OpenMP directory, along with a makefile.

#### 2.2.2 0. The basic fork-join pattern

The *omp parallel* pragma on line 22, when uncommented, tells the compiler to fork a set of threads to execute that particular line of code.

```

1  /* forkJoin.c
2   * ... illustrates the fork-join pattern
3   *      using OpenMP's parallel directive.
4   *
5   * Joel Adams, Calvin College, November 2009.
6   *
7   * Usage: ./forkJoin
8   *
9   * Exercise:
10  * - Compile & run, uncomment the pragma,
11  *   recompile & run, compare results.
12  */
13
14  #include <stdio.h>
15  #include <omp.h>
16  #include <stdlib.h>
17
18  int main(int argc, char** argv) {
19
20      printf("\nBefore...\n");
21
22      // #pragma omp parallel
23      printf("\nDuring...");
24
25      printf("\n\nAfter...\n\n");
26
27      return 0;
28  }
```

*file: patternlets/OpenMP/00.forkJoin/forkJoin.c*

### 2.2.3 1. Fork-join: setting the number of threads

Note how there is an OpenMP function for setting the number of threads to use in the next ‘fork’.

```

1  /* forkJoin2.c
2  * ... illustrates the fork-join pattern
3  *      using multiple OpenMP parallel directives,
4  *      and changing the number of threads two ways.
5  *
6  * Joel Adams, Calvin College, May 2013.
7  *
8  * Usage: ./forkJoin2
9  *
10 * Exercise:
11 * - Compile & run, compare results to source.
12 * - Predict how many threads will be used in 'Part IV'?
13 * - Uncomment 'Part IV', recompile, rerun.
14 */
15
16 #include <stdio.h>
17 #include <omp.h>
18 #include <stdlib.h>
19
20 int main(int argc, char** argv) {
21
22     printf("\nBeginning\n");
23
24     #pragma omp parallel
25     printf("\nPart I");
26
27     printf("\n\nBetween I and II...\n");
28
29     omp_set_num_threads(3);
30
31     #pragma omp parallel
32     printf("\nPart II...");
33
34     printf("\n\nBetween II and III...\n");
35
36     #pragma omp parallel num_threads(5)
37     printf("\nPart III...");
38
39     /*
40     printf("\n\nBetween III and IV...\n");
41
42     #pragma omp parallel
43     printf("\nPart IV...");
44 */
45     printf("\n\nEnd\n\n");
46
47     return 0;
48 }

```

*file patternlets/OpenMP/01.forkJoin2/forkJoin2.c*

## 2.2.4 2. Single Program, multiple data

Note how there are OpenMP functions to obtain a thread number and the total number of threads. We have one program, but multiple threads executing, each with a copy of the `id` and `num_threads` variables.

When the pragma is uncommented, note what the default number of threads is. Here the threads are forked and execute the block of code inside the curly braces on lines 22 through 26.

```

1  /* spmd.c
2  * ... illustrates the single-program-multiple-data (SPMD)
3  *    pattern using two basic OpenMP commands...
4  *
5  * Joel Adams, Calvin College, November 2009.
6  *
7  * Usage: ./spmd
8  *
9  * Exercise:
10 * - Compile & run
11 * - Uncomment pragma, recompile & run, compare results
12 */
13
14 #include <stdio.h>
15 #include <omp.h>
16 #include <stdlib.h>
17
18 int main(int argc, char** argv) {
19     printf("\n");
20
21     // #pragma omp parallel
22     {
23         int id = omp_get_thread_num();
24         int numThreads = omp_get_num_threads();
25         printf("Hello from thread %d of %d\n", id, numThreads);
26     }
27
28     printf("\n");
29     return 0;
30 }
```

file: *patternlets/OpenMP/02.spmd/spmd.c*

## 2.2.5 3. Single Program, multiple data

Here we enter the number of threads to use on the command line.

```

/* spmd2.c
 * ... illustrates the SPMD pattern in OpenMP,
 *    using the commandline arguments
 *    to control the number of threads.
 *
 * Joel Adams, Calvin College, November 2009.
 *
 * Usage: ./spmd2 [numThreads]
 *
 * Exercise:
 * - Compile & run with no commandline args
 * - Rerun with different commandline args
 */
```

```
#include <stdio.h>
#include <omp.h>
#include <stdlib.h>

int main(int argc, char** argv) {
    int id, numThreads;

    printf("\n");
    if (argc > 1) {
        omp_set_num_threads( atoi(argv[1]) );
    }

    #pragma omp parallel
    {
        id = omp_get_thread_num();
        numThreads = omp_get_num_threads();
        printf("Hello from thread %d of %d\n", id, numThreads);
    }

    printf("\n");
    return 0;
}
```

file: patternlets/OpenMP/03.spm2/spmd2.c

## 2.2.6 4. Barrier

Note what happens with and without the commented pragma on line 35.

```
/* barrier.c
 * ... illustrates the use of the OpenMP barrier command,
 *      using the commandline to control the number of threads...
 *
 * Joel Adams, Calvin College, May 2013.
 *
 * Usage: ./barrier [numThreads]
 *
 * Exercise:
 * - Compile & run several times, noting interleaving of outputs.
 * - Uncomment the barrier directive, recompile, rerun,
 *   and note the change in the outputs.
 */
```

```
#include <stdio.h>
#include <omp.h>
#include <stdlib.h>

int main(int argc, char** argv) {
    int id, numThreads;

    printf("\n");
    if (argc > 1) {
        omp_set_num_threads( atoi(argv[1]) );
    }
}
```

```

#pragma omp parallel
{
    id = omp_get_thread_num();
    numThreads = omp_get_num_threads();
    printf("Thread %d of %d is BEFORE the barrier.\n", id, numThreads);

//      #pragma omp barrier

    printf("Thread %d of %d is AFTER the barrier.\n", id, numThreads);
}

printf("\n");
return 0;
}

```

file: patternlets/OpenMP/04.barrier/barrier.c

## 2.2.7 5. Master-Worker Implementation Strategy

```

/* masterWorker.c
 * ... illustrates the master-worker pattern in OpenMP
 *
 * Joel Adams, Calvin College, November 2009.
 * Usage: ./masterWorker
 * Exercise:
 * - Compile and run as is.
 * - Uncomment #pragma directive, re-compile and re-run
 * - Compare and trace the different executions.
 */

#include <stdio.h>
#include <omp.h>

int main(int argc, char** argv) {
    int id = -1, numThreads = -1;

    printf("\n");

    #pragma omp parallel
    {
        id = omp_get_thread_num();
        numThreads = omp_get_num_threads();

        if ( id == 0 ) { // thread with ID 0 is master
            printf("Greetings from the master, # %d of %d threads\n",
                id, numThreads);
        } else { // threads with IDs > 0 are workers
            printf("Greetings from a worker, # %d of %d threads\n",
                id, numThreads);
        }
    }

    printf("\n");

    return 0;
}

```

file: patternlets/OpenMP/05.masterWorker/masterWorker.c

## 2.2.8 6. Shared Data Decomposition Pattern: blocking of threads in a parallel for loop

```
/* parallelForBlocks.c
 * ... illustrates the use of OpenMP's default parallel for loop
 *      in which threads iterate through blocks of the index range
 *      (cache-beneficial when accessing adjacent memory locations).
 *
 * Joel Adams, Calvin College, November 2009.
 * Usage: ./parallelForBlocks [numThreads]
 * Exercise
 */

#include <stdio.h>
#include <omp.h>
#include <stdlib.h>

int main(int argc, char** argv) {
    const int REPS = 8;
    int id, i;

    printf("\n");
    if (argc > 1) {
        omp_set_num_threads( atoi(argv[1]) );
    }

    #pragma omp parallel for private(id, i)
    for (i = 0; i < REPS; i++) {
        id = omp_get_thread_num();
        printf("Thread %d performed iteration %d\n",
              id, i);
    }

    printf("\n");
    return 0;
}
```

file: patternlets/OpenMP/06.parallelForLoop-blocks/parallelForBlocks.c

## 2.2.9 7. Shared Data Decomposition Pattern: striping of threads in a parallel for loop

```
/* parallelForStripes.c
 * ... illustrates how to make OpenMP map threads to
 *      parallel for-loop iterations in 'stripes' instead of blocks
 *      (use only when not accessing memory).
 *
 * Joel Adams, Calvin College, November 2009.
 *
 * Usage: ./parallelForStripes [numThreads]
 */

#include <stdio.h>
```

```

#include <omp.h>
#include <stdlib.h>

int main(int argc, char** argv) {
    const int REPS = 8;

    printf("\n");
    if (argc > 1) {
        omp_set_num_threads( atoi(argv[1]) );
    }

    #pragma omp parallel
    {
        int id = omp_get_thread_num();
        int numThreads = omp_get_num_threads();
        int i;
        for (i = id; i < REPS; i += numThreads) {
            printf("Thread %d performed iteration %d\n",
                  id, i);
        }
    }

    printf("\n");
    return 0;
}

```

file: *atternlets/OpenMP/07.parallelForLoop-stripes/parallelForStripes.c*

## 2.2.10 8. Collective Communication: Reduction

Once threads have performed independent concurrent computations, possibly on some portion of decomposed data, it is quite common to then *reduce* those individual computations into one value. In this example, an array of randomly assigned integers represents a set of shared data. All values in the array are summed together by using the OpenMP parallel for pragma with the *reduction(+:sum)* clause on the variable **sum**, which is computed in line 61.

```

1  /* reduction.c
2  * ... illustrates the OpenMP parallel-for loop's reduction clause
3  *
4  * Joel Adams, Calvin College, November 2009.
5  *
6  * Usage: ./reduction
7  * Exercise:
8  * - Compile and run. Note that correct output is produced.
9  * - Uncomment #pragma in function parallelSum()
10 *   but leave its reduction clause commented out
11 * - Recompile and rerun. Note that correct output is NOT produced.
12 * - Uncomment 'reduction(+:sum)' clause of #pragma in parallelSum()
13 * - Recompile and rerun. Note that correct output is produced again.
14 */
15
16 #include <stdio.h>
17 #include <omp.h>
18 #include <stdlib.h>
19 #include <assert.h>
20
21 void initialize(int* a, int n);
22 int sequentialSum(int* a, int n);

```



```

23 int parallelSum(int* a, int n);
24
25 #define SIZE 1000000
26
27 int main(int argc, char** argv) {
28     int array[SIZE];
29
30     initialize(array, SIZE);
31     assert( sequentialSum(array, SIZE) == parallelSum(array, SIZE) );
32     printf("Sequential and parallel functions produced the same result\n");
33
34     return 0;
35 }
36
37 /* fill array with random values */
38 void initialize(int* a, int n) {
39     int i;
40     for (i = 0; i < n; i++) {
41         a[i] = rand() % 1000;
42     }
43 }
44
45 /* sum the array sequentially */
46 int sequentialSum(int* a, int n) {
47     int sum = 0;
48     int i;
49     for (i = 0; i < n; i++) {
50         sum += a[i];
51     }
52     return sum;
53 }
54
55 /* sum the array using multiple threads */
56 int parallelSum(int* a, int n) {
57     int sum = 0;
58     int i;
59     // #pragma omp parallel for // reduction(+:sum)
60     for (i = 0; i < n; i++) {
61         sum += a[i];
62     }
63     return sum;
64 }

```

*file: patternlets/OpenMP/08.parallelForLoop-reduction/reduction.c*

## 2.2.11 9. Shared Data Pattern: Parallel-for-loop needs non-shared, private variables

```

/* private.c
 * ... illustrates why private variables are needed with OpenMP's parallel for loop
 *
 * Joel Adams, Calvin College, November 2009.
 * Usage: ./private
 * Exercise:
 * - Run, noting that the serial program produces correct results
 * - Uncomment line A, recompile/run and compare
 * - Recomment line A, uncomment line B, recompile/run and compare
 */

```

```

#include <stdio.h>
#include <omp.h>
#include <stdlib.h>

#define SIZE 100

int main(int argc, char** argv) {
    int i, j, ok = 1;
    int m[SIZE][SIZE];

    printf("\n");
    // set all array entries to 1
    // #pragma omp parallel for // A
    // #pragma omp parallel for private(j) // B
    for (i = 0; i < SIZE; i++) {
        for (j = 0; j < SIZE; j++) {
            m[i][j] = 1;
        }
    }

    // test (without using threads)
    for (i = 0; i < SIZE; i++) {
        for (j = 0; j < SIZE; j++) {
            if ( m[i][j] != 1 ) {
                printf("Element [%d,%d] not set... \n", i, j);
                ok = 0;
            }
        }
    }

    if ( ok ) {
        printf("\nAll elements correctly set to 1\n\n");
    }

    return 0;
}

```

file: patternlets/OpenMP/09.parallelForLoop-private/private.c

## 2.2.12 10. Race Condition: missing the mutual exclusion pattern

```

/* atomic.c
 * ... illustrates a race condition when multiple threads write to a shared variable
 * (and explores OpenMP private variables and atomic operations).
 *
 * Joel Adams, Calvin College, November 2009.
 * Usage: ./atomic
 * Exercise:
 * - Compile and run 10 times; note that it always produces the final balance 0
 * - To parallelize, uncomment A1+A2, recompile and rerun, compare results
 * - Try 1: recomment A1+A2, uncomment B1+B2, recompile/run, compare
 * - To fix: recomment B1+B2, uncomment A1+A2, C1+C2, recompile and rerun, compare
 */

#include<stdio.h>
#include<omp.h>

```

```

int main() {
    const int REPS = 1000000;
    int i;
    double balance = 0.0;

    printf("\nYour starting bank account balance is %.2f\n", balance);

    // simulate many deposits
    // #pragma omp parallel for // A1
    // #pragma omp parallel for private(balance) // B1
    for (i = 0; i < REPS; i++) {
        // #pragma omp atomic // C1
        balance += 10.0;
    }

    printf("\nAfter %d $10 deposits, your balance is %.2f\n",
        REPS, balance);

    // simulate the same number of withdrawals
    // #pragma omp parallel for // A2
    // #pragma omp parallel for private(balance) // B2
    for (i = 0; i < REPS; i++) {
        // #pragma omp atomic // C2
        balance -= 10.0;
    }

    // balance should be zero
    printf("\nAfter %d $10 withdrawals, your balance is %.2f\n\n",
        REPS, balance);

    return 0;
}

```

file: patternlets/OpenMP/10.mutualExclusion-atomic/atomic.c

## 2.2.13 11. Mutual Exclusion: two ways to ensure

```

/* critical.c
 * ... fixes a race condition when multiple threads write to a shared variable
 *      using the OpenMP critical directive.
 *
 * Joel Adams, Calvin College, November 2009.
 * Usage: ./critical
 * Exercise:
 * - Compile and run several times; note that it always produces the final balance 0
 * - Comment out A1+A2; recompile/run and note incorrect result
 * - To fix: uncomment B1a+B1b+B1c, B2a+B2b+B2c, recompile and rerun, compare
 */

#include<stdio.h>
#include<omp.h>

int main() {
    const int REPS = 1000000;
    int i;
    double balance = 0.0;

```

```

printf("\nYour starting bank account balance is %0.2f\n", balance);

// simulate many deposits
#pragma omp parallel for
for (i = 0; i < REPS; i++) {
    #pragma omp atomic                // A1
    #pragma omp critical              // B1a
    {                                // B1b
        balance += 10.0;
    }                                // B1c
}

printf("\nAfter %d $10 deposits, your balance is %0.2f\n",
       REPS, balance);

// simulate the same number of withdrawals
#pragma omp parallel for
for (i = 0; i < REPS; i++) {
    #pragma omp atomic                // A2
    #pragma omp critical              // B2a
    {                                // B2b
        balance -= 10.0;
    }                                // B2c
}

// balance should be zero
printf("\nAfter %d $10 withdrawals, your balance is %0.2f\n\n",
       REPS, balance);

return 0;
}

```

file: patternlets/OpenMP/11.mutualExclusion-critical/critical.c

## 2.2.14 12. Mutual Exclusion Pattern: compare performance

```

/* critical.c
 * ... compares the performance of OpenMP's critical and atomic directives
 *
 * Joel Adams, Calvin College, November 2009.
 * Usage: ./critical
 * Exercise:
 * - Compile, run, compare times for critical vs. atomic
 * - Compute how much more costly critical is than atomic, on average
 * - Create an expression setting a shared variable that atomic cannot handle (research)
 */

#include<stdio.h>
#include<omp.h>

void print(char* label, int reps, double balance, double total, double average) {
    printf("\nAfter %d $10 deposits using '%s': \n
           \n\t- balance = %0.2f, \n
           \n\t- total time = %0.12f, \n
           \n\t- average time per deposit = %0.12f\n\n",
           reps, label, balance, total, average);
}

```

```

}

int main() {
    const int REPS = 1000000;
    int i;
    double balance = 0.0,
           startTime = 0.0,
           stopTime = 0.0,
           totalTime = 0.0;

    printf("\nYour starting bank account balance is %0.2f\n", balance);

    // simulate many deposits using atomic
    startTime = omp_get_wtime();
    #pragma omp parallel for
    for (i = 0; i < REPS; i++) {
        #pragma omp atomic
        balance += 10.0;
    }
    stopTime = omp_get_wtime();
    totalTime = stopTime - startTime;
    print("atomic", REPS, balance, totalTime, totalTime/REPS);

    // simulate the same number of deposits using critical
    balance = 0;
    startTime = omp_get_wtime();
    #pragma omp parallel for
    for (i = 0; i < REPS; i++) {
        #pragma omp critical
        {
            balance += 10.0;
        }
    }
    stopTime = omp_get_wtime();
    totalTime = stopTime - startTime;
    print("critical", REPS, balance, totalTime, totalTime/REPS);

    return 0;
}

```

file: patternlets/OpenMP/12.mutualExclusion-critical2/critical2.c

## 2.2.15 13. Task Decomposition Pattern using OpenMP section directive

```

/* sections.c
 * ... illustrates the use of OpenMP's parallel section/sections directives...
 * Joel Adams, Calvin College, November 2009.
 */

#include <stdio.h>
#include <omp.h>
#include <stdlib.h>

int main(int argc, char** argv) {

    printf("\nBefore...\n\n");

```

```
#pragma omp parallel sections num_threads(4)
{
    #pragma omp section
    {
        printf("Section A performed by thread %d\n",
            omp_get_thread_num() );
    }
    #pragma omp section
    {
        printf("Section B performed by thread %d\n",
            omp_get_thread_num() );
    }
    #pragma omp section
    {
        printf("Section C performed by thread %d\n",
            omp_get_thread_num() );
    }
    #pragma omp section
    {
        printf("Section D performed by thread %d\n",
            omp_get_thread_num() );
    }
}

printf("\nAfter...\n\n");

return 0;
}
```

*file: patternlets/OpenMP/13.sections/sections.c*